

A First Look at the Triplicates & Model Fitting

Filipe Russo

March 24, 2019

Introduction

Our aim is to prepare the Proteomics dataset of the microalgae *Chrorella vulgaris* in such a way that we can construct a network of co-expressed proteins. For this objective we chose to adapt the R package WGCNA (Weighted Gene Co-expression Network Analysis), we need it to work with proteins instead of genes.

Loading & Cleaning

First, we load our original dataset `proteinGroups.txt`, store it in the variable `Proteins` and do some previous seen cleanup:

```
library(dplyr) # for data manipulation
library(stringr) # for string manipulation
library(WGCNA) # load the WGCNA package
library(tidyr) # for spread function (long to wide)

# The following setting is important for WGCNA, do not omit.
options(stringsAsFactors = FALSE)

# little nice function to check if an element IS NOT in a list
'%ni%' <- Negate('%in%')

# 1.a Loading Expression Data

Proteins = read.csv("proteinGroups.txt", sep = "\t", header=TRUE)

names(Proteins) <- str_replace(names(Proteins),
                               pattern = "MixoTP4",
                               replacement = "MixoTP6")
names(Proteins) <- str_replace(names(Proteins),
                               pattern = "HeteTP4",
                               replacement = "HeteTP10")

# we define positive identifiers for the Reverse,
# Only.identified.by.site
# and Potential.contaminant columns
rev_posids = list("+")
site_posids = list("+")
cont_posids = list("+")

clean_proteins <- Proteins %>%

# filters out rows based on posids of rev, site and cont
# filters out rows based on the protein IDs, names with CON or REV
filter(Reverse %ni% rev_posids &
       Only.identified.by.site %ni% site_posids &
```

```

    Potential.contaminant %ni% cont_posids &
    !str_detect(Majority.protein.IDs, "CON|REV")) %>%

# keeps only the ids and the LFQ columns
select(c("Majority.protein.IDs", str_subset(names(Proteins), "LFQ"))) %>%
rename(ID = Majority.protein.IDs)

```

Hierarchical Clustering

Through *Hierarchical Clustering* we will be able to cluster our triplicates and visualize sample outliers among them.

```

# removes the ID column, transposes the clean_proteins dataframe and
# then it turns the matrix back to a data.frame class object
datExprA = as.data.frame(t(clean_proteins[, -c(1:1)]))

# picks the protein ids from clean_proteins ID column and
# stores it as column names in the datExprA data.frame
names(datExprA) = clean_proteins$ID

# picks the triplicate names stored in the clean_proteins columns and
# saves them as row names from datExprA
rownames(datExprA) = names(clean_proteins)[-c(1:1)]

# 1.b Checking data for excessive missing values and
# Identification of outlier microarray samples

gsg = goodSamplesGenes(datExprA, verbose = 3);

## Flagging genes and samples with too many missing values...
## ..step 1
gsg$allOK

## [1] TRUE
if (!gsg$allOK)
{
  # Optionally, print the gene and sample names that were removed:
  if (sum(!gsg$goodGenes)>0)
    printFlush(paste("Removing proteins:", paste(names(datExprA)[!gsg$goodGenes],
                                                    collapse = ", ")));
  if (sum(!gsg$goodSamples)>0)
    printFlush(paste("Removing samples:", paste(rownames(datExprA)[!gsg$goodSamples],
                                                    collapse = ", ")));

  # Remove the offending genes and samples from the data:
  datExprA = datExprA[gsg$goodSamples, gsg$goodGenes]
}

# Clustering dendrogram of samples based on their Euclidean distance
sampleTree1 = hclust(dist(datExprA), method = "average");

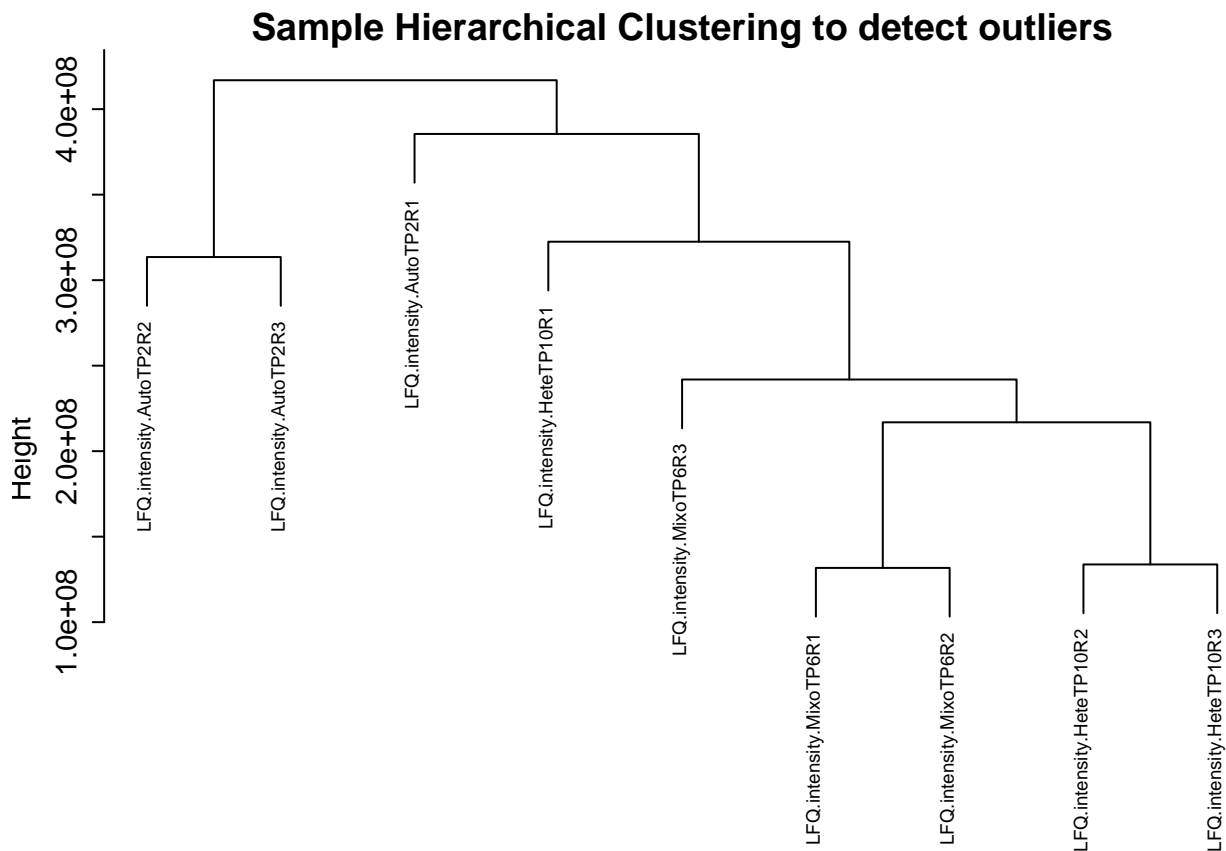
# pdf(file = "sampleClustering.pdf", width = 12, height = 9);

```

```

par(cex = 0.6);
par(mar = c(0,4,2,0))
plot(sampleTree1,
     main = "Sample Hierarchical Clustering to detect outliers",
     sub = "",
     xlab = "",
     cex.lab = 1.5,
     cex.axis = 1.5,
     cex.main = 2)

```



It seems the sample HeteTP10R1 is an outlier. We suppose this happened due to the zeros found in some LFQ intensities, some represent true zeros and others represent missing values. We want the samples to be clustered in 3 different clusters: Auto, Mixo, Hete; let's see how the `cutree` function handles it.

```
cutree(sampleTree1, k = 3)
```

```

## LFQ.intensity.AutoTP2R1 LFQ.intensity.AutoTP2R2 LFQ.intensity.AutoTP2R3
##                          1                      2                      2
## LFQ.intensity.HeteTP10R1 LFQ.intensity.HeteTP10R2 LFQ.intensity.HeteTP10R3
##                          3                      3                      3
## LFQ.intensity.MixoTP6R1  LFQ.intensity.MixoTP6R2  LFQ.intensity.MixoTP6R3
##                          3                      3                      3

```

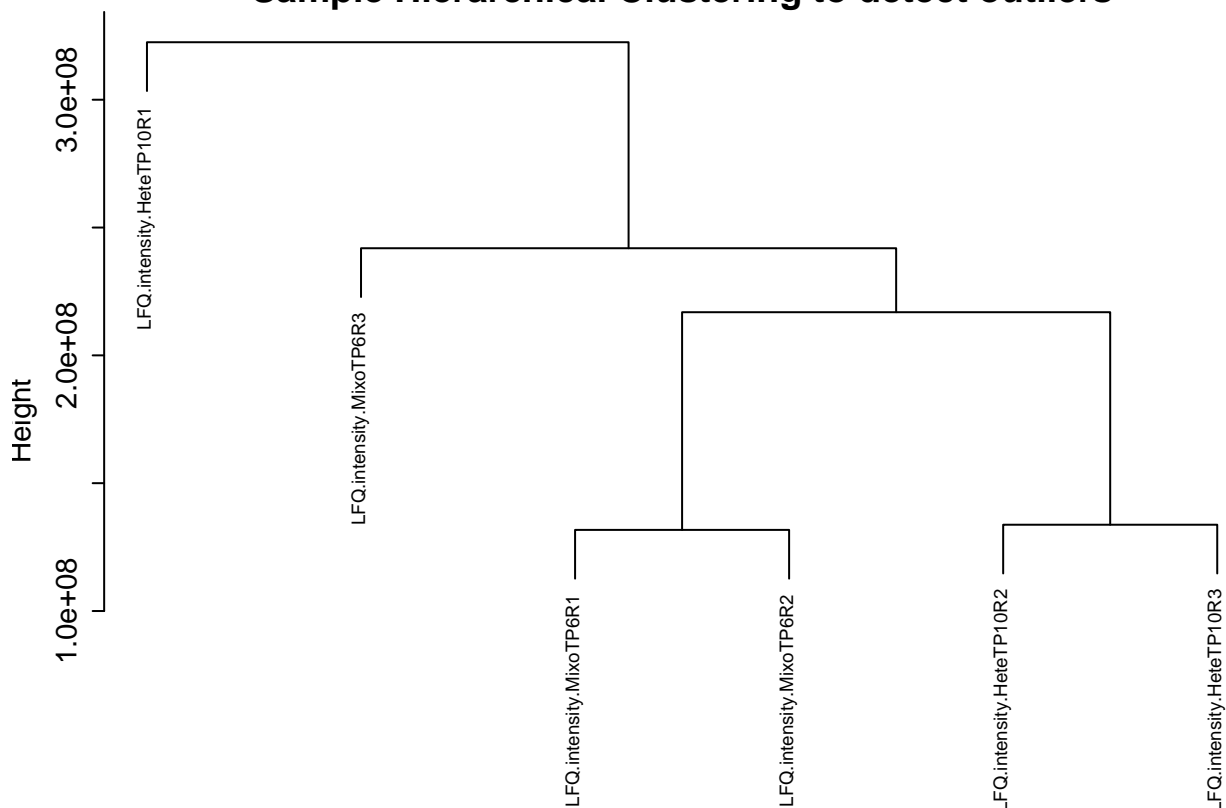
The result suggests the sample AutoTP2R1 is an outlier among the AutoTP2 replicates and the growth conditions Heterotrophic and Mixotrophic seem to be more similar to each other than to the Autotrophic one.

Let's remove the Autotrophic growth condition from the data.frame and see how `hclust` and `cutree` handle the samples:

```
# Clustering dendrogram of samples based on their Euclidean distance
sampleTree2 = hclust(dist(datExprA[-c(1:3), ]), method = "average");

# pdf(file = "sampleClustering.pdf", width = 12, height = 9);
par(cex = 0.6);
par(mar = c(0,4,2,0))
plot(sampleTree2,
     main = "Sample Hierarchical Clustering to detect outliers",
     sub = "",
     xlab = "",
     cex.lab = 1.5,
     cex.axis = 1.5,
     cex.main = 2)
```

Sample Hierarchical Clustering to detect outliers



```
cutree(sampleTree2, k = 2)
```

```
## LFQ.intensity.HeteTP10R1 LFQ.intensity.HeteTP10R2 LFQ.intensity.HeteTP10R3
##           1                2                2
## LFQ.intensity.MixoTP6R1 LFQ.intensity.MixoTP6R2 LFQ.intensity.MixoTP6R3
##           2                2                2
```

As expected the HeteTP10R1 sample seems to be an outlier, let's remove it and redo the process:

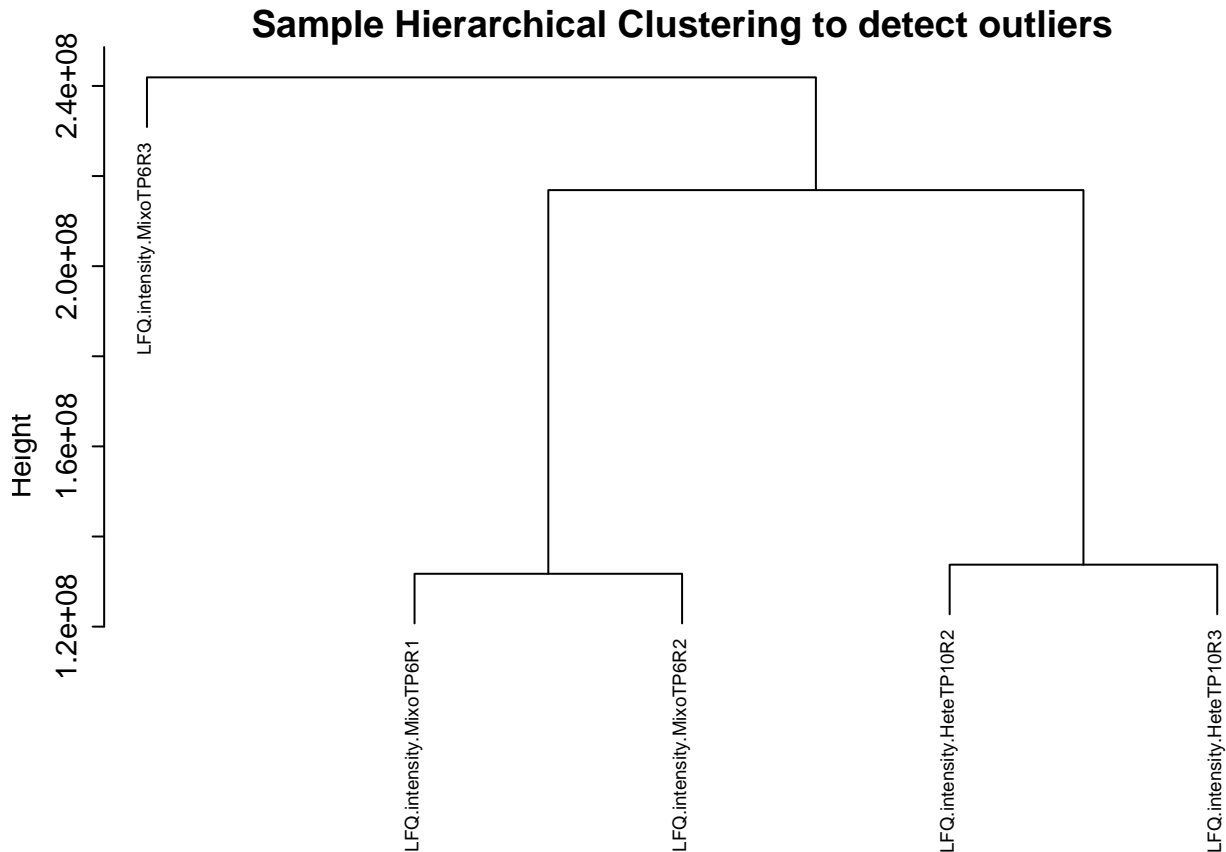
```
# Clustering dendrogram of samples based on their Euclidean distance
sampleTree3 = hclust(dist(datExprA[-c(1:4), ]), method = "average");

# pdf(file = "sampleClustering.pdf", width = 12, height = 9);
par(cex = 0.6);
```

```

par(mar = c(0,4,2,0))
plot(sampleTree3,
     main = "Sample Hierarchical Clustering to detect outliers",
     sub = "",
     xlab = "",
     cex.lab = 1.5,
     cex.axis = 1.5,
     cex.main = 2)

```



```

cutree(sampleTree3, k = 2)

```

```

## LFQ.intensity.HeteTP10R2 LFQ.intensity.HeteTP10R3 LFQ.intensity.MixoTP6R1
##           1           1           1
## LFQ.intensity.MixoTP6R2 LFQ.intensity.MixoTP6R3
##           1           2

```

Now, the sample MixoTP6R3 seems to be the outlier. Let's remove it, take one last look at this slicing process and move on with the analysis:

```


# Clustering dendrogram of samples based on their Euclidean distance
sampleTree4 = hclust(dist(datExprA[-c(1:4, 9), ]), method = "average");

# pdf(file = "sampleClustering.pdf", width = 12, height = 9);
par(cex = 0.6);
par(mar = c(0,4,2,0))
plot(sampleTree4,
     main = "Sample Hierarchical Clustering to detect outliers",
     sub = "",

```

```
xlab = "",
cex.lab = 1.5,
cex.axis = 1.5,
cex.main = 2)
```

Sample Hierarchical Clustering to detect outliers



```
cutree(sampleTree4, k = 2)
```

```
## LFQ.intensity.HeteTP10R2 LFQ.intensity.HeteTP10R3 LFQ.intensity.MixoTP6R1
##                      1                      1                      2
## LFQ.intensity.MixoTP6R2
##                      2
```

The last result seemed somewhat better than the previous ones, since we could in fact properly cluster the samples.

We will try to improve these results by taking the median value among the triplicates as the best estimative of the true LFQIntensity value.

```
# we already took the median value in the previous report,
# so now we only load the data and redo the clustering
```

```
valid_proteins_long <- read.csv("proteins.csv", sep = ",", header = TRUE)
valid_proteins_wide <- valid_proteins_long %>% select(-Time) %>%
  spread(Growth, MedLFQ)
```

```
# removes the ID column, transposes the valid_proteins_wide dataframe and
# then it turns the matrix back to a data.frame class object
datExprA2 = as.data.frame(t(valid_proteins_wide[, -c(1:1)]))
```

```

# picks the protein ids from valid_proteins_wide ID column and
# stores it as column names in the datExprA2 data.frame
names(datExprA2) = valid_proteins_wide$ID


# picks the triplicate names stored in the valid_proteins_wide columns and
# saves them as row names from datExprA2
rownames(datExprA2) = names(valid_proteins_wide)[-c(1:1)]

# Clustering dendrogram of samples based on their Euclidean distance
sampleTree5 = hclust(dist(datExprA2), method = "average");

# pdf(file = "sampleClustering.pdf", width = 12, height = 9);
par(cex = 0.6);
par(mar = c(0,4,2,0))
plot(sampleTree5,
      main = "Sample Hierarchical Clustering to detect outliers",
      sub = "",
      xlab = "",
      cex.lab = 1.5,
      cex.axis = 1.5,
      cex.main = 2)

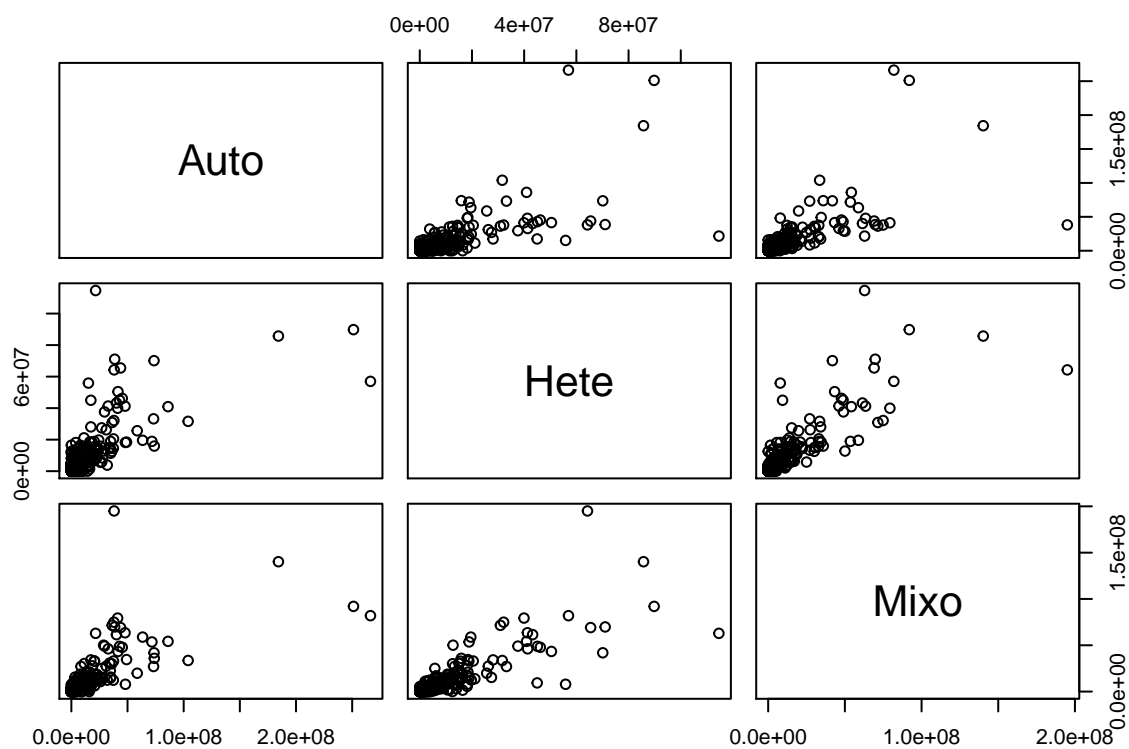
```

Sample Hierarchical Clustering to detect outliers

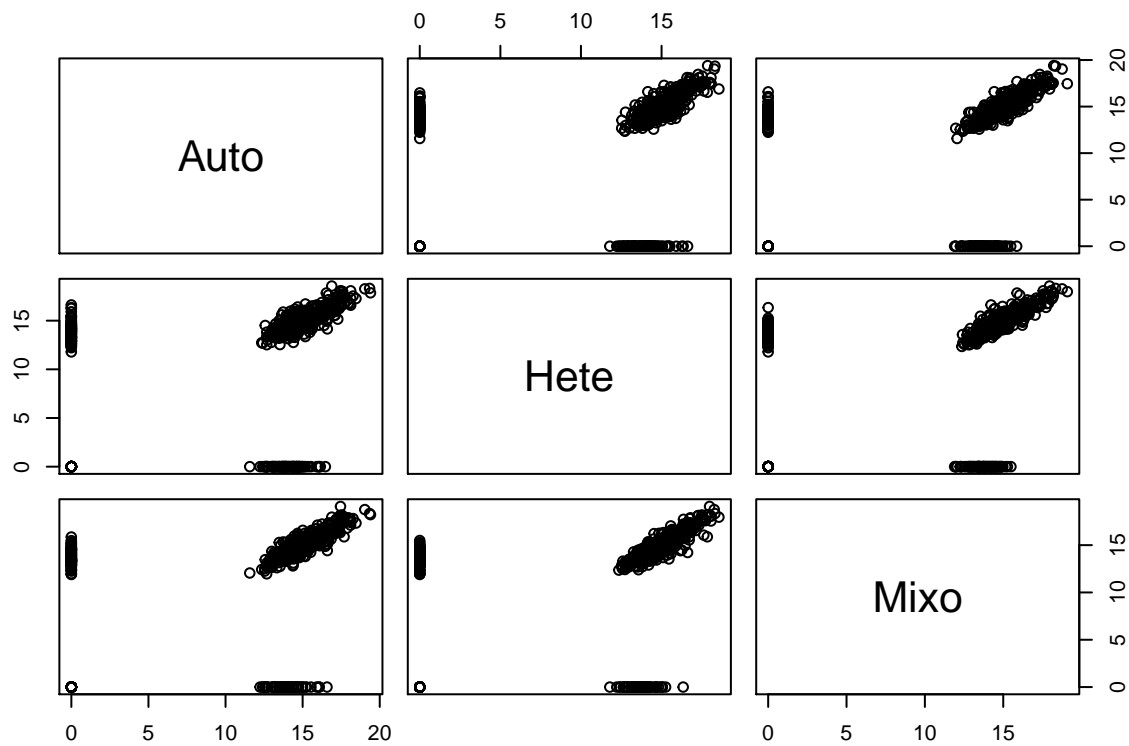


The plot raises the question: Are the Mixotrophic and Heterotrophic growth conditions more similar to each other than to the Autotrophic one?

```
pairs(valid_proteins_wide[, -1])
```



```
pairs(log(valid_proteins_wide[, -1] + 1))
```



Scale Free Topology Model Fit

The WGCNA package requires us to pick a soft threshold power to estimate both *model fit* and *mean connectivity*, so we can continue with the network construction.

First, we will work with the `datExprA` data.frame devired from the `clean_proteins` data.frame, which is a cleaned version from the original Proteins dataset.

```
# 2.a Automatic network construction and module detection
```

```
# Choose a set of soft-thresholding powers
```

```
powers = c(c(1:10), seq(from = 12, to = 40, by = 2))
```

```
# Call the network topology analysis function
```

```
sft = pickSoftThreshold(datExprA, powerVector = powers, verbose = 5)
```

```
## pickSoftThreshold: will use block size 757.
```

```
## pickSoftThreshold: calculating connectivity for given powers...
```

```
## ..working on genes 1 through 757 of 757
```

```
## Warning: executing %dopar% sequentially: no parallel backend registered
```

##	Power	SFT.R.sq	slope	truncated.R.sq	mean.k.	median.k.	max.k.
## 1	1	0.111000	1.8700	0.7800	251.00	255.000	334.0
## 2	2	0.028500	0.4610	0.8300	122.00	122.000	191.0
## 3	3	0.003800	0.1110	0.9180	71.60	69.400	124.0
## 4	4	0.000919	-0.0454	0.8880	47.60	47.900	88.4
## 5	5	0.009460	-0.1200	0.8220	34.40	34.600	66.6
## 6	6	0.000677	-0.0227	0.7630	26.30	25.600	52.1
## 7	7	0.040800	-0.1530	0.4600	21.00	19.900	44.0
## 8	8	0.263000	-0.4830	0.4940	17.40	15.700	41.0
## 9	9	0.342000	-0.9660	0.3260	14.70	12.800	40.5
## 10	10	0.182000	-2.9500	0.0688	12.80	10.400	40.1
## 11	12	0.258000	-4.4300	0.1390	10.10	7.470	39.7
## 12	14	0.284000	-5.1500	0.0931	8.49	5.380	39.5
## 13	16	0.288000	-5.0500	0.0988	7.37	4.040	39.3
## 14	18	0.292000	-4.9400	0.1040	6.59	3.130	39.2
## 15	20	0.295000	-4.8000	0.1090	6.02	2.430	39.2
## 16	22	0.282000	-4.9000	0.0789	5.59	1.910	39.1
## 17	24	0.283000	-4.7800	0.0798	5.26	1.520	39.1
## 18	26	0.282000	-4.6600	0.0782	5.00	1.250	39.1
## 19	28	0.283000	-4.5100	0.0789	4.79	1.020	39.0
## 20	30	0.276000	-4.2900	0.0718	4.62	0.859	39.0
## 21	32	0.214000	-3.8700	0.0247	4.48	0.729	39.0
## 22	34	0.282000	-4.1600	0.0784	4.36	0.633	39.0
## 23	36	0.282000	-4.0700	0.0790	4.27	0.571	39.0
## 24	38	0.282000	-3.9900	0.0792	4.18	0.494	39.0
## 25	40	0.282000	-3.9300	0.0788	4.11	0.423	39.0

```
# Plot the results:
```

```
cex1 = 0.9;
```

```
# We try to grab an automated power estimate from the sft list
```

```
sft$powerEstimate
```

```
## [1] NA
```

```

# sft$powerEstimate is the lowest power for which the scale free topology fit
# R^2 exceeds RsquaredCut. If R^2 is below RsquaredCut for all powers, NA is returned.

# Since the function pickSoftThreshold didn't give us a value in
# the sft$powerEstimate, we have to find by ourselves a value
# that is the lowest power for which the scale-free topology fit index curve
# flattens out upon reaching a high value

model_fit <- -sign(sft$fitIndices[ , "slope"]) * sft$fitIndices[ , "SFT.R.sq"]
mean_con <- sft$fitIndices[ , "mean.k."]


# Scale-free topology fit index as a function of the soft-thresholding power
plot(powers,
     model_fit,
     xlab = "Soft Threshold (power)",
     ylab = "Scale Free Topology Model Fit, signed R^2",
     type = "n",
     main = paste("Scale independence"),
     yaxt = "n")

text(powers,
     model_fit,
     labels = powers,
     cex = cex1,
     col = "blue")

axis(2, at = seq(-0.1, 0.2, 0.1))

# Model Fit value corresponding to power 16
abline(h = model_fit[13], col = "red")
axis(2, at = 0.2878373, labels = c(0.28), col = "red")

```




```
# Mean connectivity as a function of the soft-thresholding power
plot(powers,
     mean_con,
     xlab = "Soft Threshold (power)",
     ylab = "Mean Connectivity",
     type = "n",
     main = paste("Mean connectivity"),
     yaxt = "n")

text(powers,
     mean_con,
     labels = powers,
     cex = cex1,
     col = "blue")

axis(2, at = seq(50, 250, 50))

# Mean Connectivity value corresponding to power 16
abline(h = mean_con[13], col = "red")
axis(2, at = 7.370472, labels = c(7.37), col = "red")
```



As can be seen in the two previous plots we have chosen by visual analysis the Soft Threshold Power 16. The Power 16 gives us a 0.2878373 value for the Model Fit Index.

Now, we will do the same for the `datExprA2` data.frame which is based on the median value among triplicates. Let's take a look:

```
# 2.a Automatic network construction and module detection

# Choose a set of soft-thresholding powers
powers = c(c(1:10), seq(from = 12, to = 40, by = 2))

# Call the network topology analysis function
sft2 = pickSoftThreshold(datExprA2, powerVector = powers, verbose = 5)
```

```
## pickSoftThreshold: will use block size 600.
## pickSoftThreshold: calculating connectivity for given powers...
## ..working on genes 1 through 600 of 600
## Power SFT.R.sq slope truncated.R.sq mean.k. median.k. max.k.
## 1 1 0.1800 4.780 0.816 395.0 410.0 433
## 2 2 0.2390 2.640 0.753 313.0 328.0 364
## 3 3 0.2390 2.050 0.625 268.0 282.0 321
## 4 4 0.2230 1.860 0.776 238.0 253.0 289
## 5 5 0.1700 1.770 0.818 217.0 233.0 264
## 6 6 0.1470 1.750 0.879 201.0 219.0 244
## 7 7 0.1050 1.570 0.894 189.0 207.0 229
## 8 8 0.0976 1.550 0.849 179.0 198.0 217
## 9 9 0.0684 1.280 0.785 170.0 191.0 208
## 10 10 0.0287 0.769 0.687 163.0 184.0 201
## 11 12 0.0787 1.160 0.590 151.0 170.0 189
## 12 14 0.0775 1.020 0.497 141.0 158.0 180
## 13 16 0.1310 1.190 0.459 133.0 148.0 172
```

## 14	18	0.1350	1.190	0.386	127.0	139.0	166
## 15	20	0.1480	1.020	0.286	121.0	131.0	160
## 16	22	0.2210	1.000	0.357	117.0	125.0	155
## 17	24	0.2370	0.954	0.360	112.0	119.0	151
## 18	26	0.2650	0.907	0.378	109.0	114.0	148
## 19	28	0.2470	0.864	0.332	105.0	109.0	144
## 20	30	0.4680	0.898	0.600	102.0	105.0	141
## 21	32	0.4070	0.872	0.492	99.6	101.0	138
## 22	34	0.3860	0.845	0.460	97.1	97.8	136
## 23	36	0.3450	0.790	0.394	94.8	94.9	134
## 24	38	0.2920	0.744	0.313	92.6	92.3	132
## 25	40	0.1890	0.612	0.125	90.7	89.7	130

```
# Plot the results:
```

```
cex1 = 0.9;
```

```
# We try to grab an automated power estimate from the sft2 list
sft2$powerEstimate
```

```
## [1] NA
```

```
# sft2$powerEstimate is the lowest power for which the scale free topology fit
# R^2 exceeds RsquaredCut. If R^2 is below RsquaredCut for all powers, NA is returned.
```

```
# Since the function pickSoftThreshold didn't give us a value in
# the sft2$powerEstimate, we have to find by ourselves a value
# that is the lowest power for which the scale-free topology fit index curve
# flattens out upon reaching a high value
```

```
model_fit2 <- -sign(sft2$fitIndices[, "slope"]) * sft2$fitIndices[, "SFT.R.sq"]
mean_con2 <- sft2$fitIndices[, "mean.k."]
```

```
# Scale-free topology fit index as a function of the soft-thresholding power
```


```
plot(powers,
     model_fit2,
     xlab = "Soft Threshold (power)",
     ylab = "Scale Free Topology Model Fit, signed R^2",
     type = "n",
     main = paste("Scale independence"),
     yaxt = "n")
```

```
text(powers,
     model_fit2,
     labels = powers,
     cex = cex1,
     col = "blue")
```

```
#axis(2, at = seq(-0.1, 0.2, 0.1))
```

```
# Model Fit value corresponding to power 16
```

```
abline(h = model_fit2[13], col = "red")
axis(2, at = -0.1312341, labels = c(-0.13), col = "red")
```




```
# Mean connectivity as a function of the soft-thresholding power
plot(powers,
     mean_con2,
     xlab = "Soft Threshold (power)",
     ylab = "Mean Connectivity",
     type = "n",
     main = paste("Mean connectivity"),
     yaxt = "n")

text(powers,
     mean_con2,
     labels = powers,
     cex = cex1,
     col = "blue")

axis(2, at = c(200))

# Mean Connectivity value corresponding to power 16
abline(h = mean_con2[13], col = "red")
axis(2, at = 133.4083, labels = c(133), col = "red")
```

Mean connectivity



Since we couldn't find a good soft-thresholding power for the `datExprA2` data.frame we decided to use the power 16 found for the `datExprA` data.frame. Unfortunately, it gave us a worse fit both graphically and numerically, the Model Fit Index this time was -0.1312341. The results on both data frames were very much influenced by the minimum requirements of `WGCNA` package, according to the `WGCNA` FAQ one needs at least 15 samples to find results that are not too noisy, it also recommends at least 20 samples for better results.

Saving our Expression Data

Finally we save our data frames `datExprA` and `datExprA2` for further analysis.

```
write.csv(datExprA, file = "datExprA.csv", row.names = TRUE)
datExprA <- read.csv("datExprA.csv", sep = ",", header = TRUE)
rownames(datExprA) = datExprA$X
datExprA <- datExprA[ , -c(1)]

write.csv(datExprA2, file = "datExprA2.csv", row.names = TRUE)
datExprA2 <- read.csv("datExprA2.csv", sep = ",", header = TRUE)
rownames(datExprA2) = datExprA2$X
datExprA2 <- datExprA2[ , -c(1)]
```